

Approximating the optimal threshold for an
abstaining classifier based on a reward function
with regression

BACHELOR THESIS

Jonas Fassbender

jonas@fassbender.dev
11117674

In the course of studies
COMPUTER SCIENCE

For the degree of
BACHELOR OF SCIENCE

Technical University of Cologne
Faculty of Computer Science and Engineering

First supervisor: Prof. Dr. Heinrich Klocke
Technical University of Cologne

Second supervisor: Prof. Dr. Fotios Giannakopoulos
Technical University of Cologne

Overath, July 2019

1. Introduction

An abstaining classifier (see e.g. Vanderlooy et al., 2009)—also called a classifier with reject option (see e.g. Fischer et al., 2016)—is a kind of confidence predictor. It can refuse from making a prediction if its confidence in the prediction is not high enough. High enough, in this context, means that the confidence is greater than a certain—hopefully optimal—threshold. Optimality is dependent on a performance metric set beforehand.

This thesis introduces a new kind of method for approximating the optimal threshold based on a reward function—better known from reinforcement learning than from the supervised learning setting (see e.g. Sutton and Barto, 2018, Chapter 1). The method treats the reward function as unknown, making it a very general approach and giving quite the amount of freedom in designing the reward function.

In supervised learning the concept that is closest to a reward function is a cost function and many abstract types of cost in supervised learning are known (see Turney, 2002).

Probably today’s most used methods for obtaining the optimal threshold for reducing the expected cost of an abstaining classifier are based on the receiver operating characteristic (ROC) rule (see Tortorella, 2000; Pietraszek, 2005; Vanderlooy et al., 2009; Guan et al., 2018).

The method presented in this thesis is more flexible than the methods based on the ROC rule and can—depending on the context of the classification problem—produce results better interpretable than results from a cost setting (see Chapter 2). Also it is more natural with multi-class classification problems than the methods based on the ROC rule, all assuming binary classification problems, wherefore the classifiers generated by these methods must be transformed to multi-class classifiers for non-binary problems.

On the other hand the presented method can suffer from its very general approach and only produces approximations. This can result in non-optimal and unstable thresholds.

This thesis first presents a motivational example. In Chapter 3 the proposed method is presented. After that experiments on data sets from the UCI machine learning repository (see Dua and Graff, 2017) are discussed. At last further research ideas are listed and a conclusion is drawn.

2. Motivational example

This chapter will point out the usefulness of abstaining classifiers in real world application domains where reliability is key. It will show an example why the reward setting can improve readability in some domains. First another example, for which the cost setting—more commonly used in supervised learning—comes more natural is given and the differences are discussed.

Abstaining classifiers—compared to typical classifiers, which classify every prediction, maybe even without a confidence value in it (then called a bare prediction)—can be easily integrated into and enhance processes where they partially replace some of the decision

making, since they can delegate the abstained predictions back to the underlying process. The use of abstaining classifiers in domains where reliability—in regard to prediction errors—is important, has an interesting aspect in giving reliability while still being able to decrease work, cost, etc. to some degree. This is a valuable property if there does not exist a typical classifier good enough to fully replace the underlying process.

Many real world application domains for abstaining classifiers can express a cost function associated to the decisions about predicting and abstaining of the classifier—which then chooses the threshold with which it produces the least amount of cost, therefore minimizing the cost of introducing the abstaining classifier to the process.

For example, the real world application domain could be a facial recognition system at a company which regulates which employee can enter a trust zone and which can not. The process which should be enhanced with the facial recognition system is a manual process where the employee has to fill out a form in order to receive a key which opens the trust zone.

In this example, the costs of miss-classifying an unauthorized person as authorized can be huge for the company while abstaining or classifying an authorized employee as unauthorized produces quite low costs—the authorized employee just has to start the manual process, which should be replaced by the facial recognition system.

On the other hand, for some real world application domains a reward function based on which the abstaining classifier chooses the threshold by maximizing the reward—rather than minimizing the cost—comes more natural.

Such a domain would be the finance industry, where we often can associate a certain amount of money an abstaining classifier can produce or save by supporting the decision making of an underlying process.

An example for such a process would be the process of a bank for granting a consumer credit. The bank requests information about the consumer from a credit bureau in order to assess the consumer’s credit default risk. Now the bank wants to predict the consumer’s credit default risk based on information the bank has about the consumer. If the credit default risk is very high or very low the bank can save money not making a request to the credit bureau for this consumer. The optimal threshold for the abstaining classifier making the prediction about the credit default risk can easily be expressed by a reward function. Every correct decision saves the bank the money the request to the credit bureau costs. Every miss-classification costs the bank either the amount of money it would gain by granting the credit, or the money it loses by giving a credit to somebody that does not pay the rates. Abstention cost is the cost of making a request to the credit bureau.

Using a reward function—like in the example above—instead of a cost function has an advantage in readability. One can easily assess the gain of introducing the abstaining classifier to the process. Is the reward generated by the abstaining classifier higher than zero, the process is enhanced by the abstaining classifier. Otherwise the abstaining classifier would produce more cost than it would save and it is not valuable for the bank to introduce it to its process of assessing a consumer’s credit default risk.

3. Proposed method based on reward

Let \mathbf{X} be our observation space and \mathbf{Y} our label space. $|\mathbf{Y}| < \infty$ since only classification is discussed. Let \mathbf{Z} be the cartesian product of \mathbf{X} and \mathbf{Y} : $\mathbf{Z} := \mathbf{X} \times \mathbf{Y}$. \mathbf{Z} is called the example space. Let an example z_i from \mathbf{Z} be: $z_i := (x_i, y_i)$; $z_i \in \mathbf{Z}$. A data set¹ containing examples z_1, \dots, z_n is annotated as $\{z_1, \dots, z_n\}$.

3.1 Scoring classifiers

A classical machine learning predictor—in the previous chapter called a typical classifier—can be represented by a function

$$D : \mathbf{Z}^* \times \mathbf{X} \rightarrow \mathbf{Y}. \quad (1)$$

Its first argument being a data set with an arbitrary length the classifier is trained on, while the second is an observation which should be predicted (mapped to a label from \mathbf{Y}).

Let $D_{\{z_1, \dots, z_n\}}$ be a classical machine learning predictor trained on the data set $\{z_1, \dots, z_n\}$ and let $D_{\{z_1, \dots, z_n\}}(x)$ be equivalent to (1) with the first argument being $\{z_1, \dots, z_n\}$.

The proposed method relies on scoring classifiers. A scoring classifier does not return just a label but instead returns some score for each label from our label space. The only constraint on the scores is that higher scores are better than lower. A score could be a probability or just an uncalibrated confidence value (see Vanderlooy et al., 2009).

Let S be a scoring classifier:

$$S : \mathbf{Z}^* \times \mathbf{X} \rightarrow (\mathbf{Y} \rightarrow \mathbb{R}). \quad (2)$$

S takes the same arguments as (1) but instead of producing bare predictions it returns a function which maps every label from the label space to a score determined by S .

The method proposed is only interested in the highest score and the associated label. For that two functions k and v are defined:

$$\begin{aligned} k(S_{\{z_1, \dots, z_n\}}, x) &= \arg \max_{y \in \mathbf{Y}} S_{\{z_1, \dots, z_n\}}(x)(y) \\ v(S_{\{z_1, \dots, z_n\}}, x) &= \max_{y \in \mathbf{Y}} S_{\{z_1, \dots, z_n\}}(x)(y). \end{aligned}$$

The composition kv of k and v returns the tuple with the label mapped to the highest score:

$$kv(S_{\{z_1, \dots, z_n\}}, x) = (k(S_{\{z_1, \dots, z_n\}}, x), v(S_{\{z_1, \dots, z_n\}}, x)). \quad (3)$$

1. not an actual set but a multi-set since it can contain the same element more often than one time.

3.2 Abstaining classifiers

An abstaining classifier A can be defined as a similar function as (1), with the only difference being the return value:

$$A : \mathbf{Z} \times \mathbf{X} \rightarrow \mathbf{Y} \cup \{\perp\}$$

A can return a label from \mathbf{Y} , but also \perp , indicating that A would like to abstain from making a prediction.

Let $\mathbf{S}_{\mathbf{Z}}$ be the set of all scoring classifiers defined like (2) on the example set \mathbf{Z} . The proposed method is interested in transforming a scoring classifier $S \in \mathbf{S}_{\mathbf{Z}}$ to an abstaining classifier A . In order to do that a threshold $T \in \mathbb{R}$ is defined and A can be represented as a composition of S and T . Let $S_{\langle z_1, \dots, z_n \rangle}$ be a scoring classifier, T a threshold and x an observation to be predicted. The abstaining classifier A composed of $S_{\langle z_1, \dots, z_n \rangle}$ and T predicts x as follows:

$$A(\langle z_1, \dots, z_n \rangle, x) = \begin{cases} k(S_{\langle z_1, \dots, z_n \rangle}, x) & \text{if } v(S_{\langle z_1, \dots, z_n \rangle}, x) > T \\ \perp & \text{if } v(S_{\langle z_1, \dots, z_n \rangle}, x) \leq T \end{cases}.$$

This representation of A is rather unconventional and is one reason the proposed method is unstable.

Using a single threshold for all labels is a strong constraint to put onto the scoring classifier, because it must be invariant to the label distribution. Imagine a binary problem $\mathbf{Y} = \{P, N\}$, where P is the positive label and N the negative. Now, P makes up 90 percent of all examples and the scoring classifier is not invariant to the label distribution. This could lead the classifier to produce higher scores for P than for N which could result in an abstaining classifier that does not predict on any example with the label N , even though with such a distribution predicting N would probably be more interesting.

ROC based and other methods for generating abstaining classifiers address this problem by using abstention windows instead of a single threshold (see Friedel et al., 2006).

The margin $m : \mathbf{Y} \times \mathbb{R} \rightarrow (-1, 1)$ is a function that combines the label with the confidence value and returns a number in the interval of $(-1, 1)$. The closer the return value of m is to the edges, the more confident the scoring classifier is, whereby -1 means perfectly confident the label is N and 1 means perfectly confident the label is P (see Friedel et al., 2006).

In Guan et al. (2018) a similar method is described, constraining the output of the margin m not on $(-1, 1)$ but instead using only the likelihood of an observation x having the positive label P ($m : \mathbb{R} \rightarrow (0, 1)$).

Both Friedel et al. (2006) and Guan et al. (2018) define an abstention window a as a tuple $a := (t_1, t_2); t_1 < t_2$ with two thresholds. An abstaining classifier with an abstention

window instead of a threshold predicts an observation x as:

$$A(z_1, \dots, z_n, x) = \begin{cases} P & \text{if } m(kv(S_{z_1, \dots, z_n}, x)) > t_2 \\ \perp & \text{if } t_1 \leq m(kv(S_{z_1, \dots, z_n}, x)) \leq t_2 \\ N & \text{if } m(kv(S_{z_1, \dots, z_n}, x)) < t_1 \end{cases}.$$

This addresses the problem of using a single threshold T for predictions on both labels from \mathbf{Y} . The constraint of abstention windows is that they are only defined on binary problems, while an arbitrary number of labels can be predicted with a single threshold.

This thesis does not address the problem of using a single threshold, but a possible solution is given in Chapter 5.

3.3 Abstaining classifiers from reward

The novel approach of this thesis is using a system based on reward which is maximized rather than cost that is minimized in order to determine the optimal threshold for abstention. Like stated in Chapter 1 using a reward function—like used in reinforcement learning—in a supervised learning setting is rather uncommon. In Chapter 1 and Chapter 2 the reasons why using reward instead of cost are given.

An aspect of cost—which makes it less flexible than reward—is that it is only defined on \mathbb{R}^+ , while reward is defined on \mathbb{R} . Reward combines cost with gain.

Let ρ be a reward function:

$$\rho : \mathbf{Y}^* \times \mathbf{Y}^* \rightarrow \mathbb{R}^*.$$

ρ takes two arbitrary long vectors with labels from \mathbf{Y} . The first vector contains the true labels of some sequence of examples, the second contains the predicted labels from a scoring classifier S for the same sequence. The sequence of examples is ordered after the scores produced by S in descending order. ρ returns a vector of rewards for each tuple of true label and predicted label.

The reward function is basically treated as a black box function, the only knowledge we have is, whether ρ produces single-step reward or accumulated reward values.

Treating the reward function this way makes it much more flexible than the cost setting which uses cost matrices (see Fischer et al., 2016). A cost matrix C for a binary abstaining classifier is defined as

$$C := \begin{pmatrix} C(P, P) & C(P, N) & C(P, \perp) \\ C(N, P) & C(N, N) & C(N, \perp) \end{pmatrix}$$

and is basically the inverse of a single step reward function—with the difference that $C(P, P)$ and $C(N, N)$ normally do not have a gain associated to them, because then the cost matrix would not be true to its cost setting. A cell of C would provide a gain if its value is smaller than zero.

A reward function that accumulates provides an even more flexible setting than single step reward—which is only dependent on one example’s true and predicted label—because it can introduce the concept of state (see Sutton and Barto, 2018, Chapter 1).

For example, our classifier could be a bettor betting on the outcome of a card game. It starts with a certain amount of money and always bets two thirds of its amount. Every example is one match and it is possible to derive a certainty measure based on some information about the match. The reward is the amount of money the classifier wins or loses. It gains a certain amount—depending on how much money the classifier owns after the last match it has bet on—if it decides to bet on the current match and does so correctly or loses two thirds of its reward if the classifier was wrong. A reward function like this is not stateless like a single step reward function and is a classical alternative to it in reinforcement learning (see Sutton and Barto, 2018, Chapter 1).

The idea now is to find a threshold with which the expected reward is maximized.

3.4 Equivalences to reinforcement learning

4. Experiments

5. Further research

6. Conclusion

Appendix

A. Plots

References

- Dheeru Dua and Casey Graff. UCI machine learning repository, 2017. URL <http://archive.ics.uci.edu/ml>.
- Lydia Fischer, Barbara Hammer, and Heiko Wersing. Optimal local rejection for classifiers. *Neurocomputing*, 214:445 – 457, 2016. ISSN 0925-2312. doi: <https://doi.org/10.1016/j.neucom.2016.06.038>. URL <http://www.sciencedirect.com/science/article/pii/S0925231216306762>.
- Caroline C. Friedel, Ulrich Rückert, and Stefan Kramer. Cost curves for abstaining classifiers. In *Proceedings of the ICML 2006 workshop on ROC Analysis in Machine Learning*, 2006.
- Hongjiao Guan, Yingtao Zhang, Heng-Da Cheng, and Xianglong Tang. Abstaining classification when error costs are unequal and unknown. *CoRR*, abs/1806.03445, 2018. URL <http://arxiv.org/abs/1806.03445>.

Algorithm 1 : k-fold method for determining the threshold for an abstaining classifier based on reward

Input:

S : a scoring classifier,
 ρ : a reward function,
data set: $\{z_1, \dots, z_n\}$,
 k : the amount of partitions,
 R : a regressor (optionally)

Output:

T: threshold

- 1: predicted points $:= \{\}$
- 2: split data set into k roughly equal sized partitions $\text{split}_1, \dots, \text{split}_k$
- 3: **for all** $\text{split}_i, i = 1, \dots, k$ **do**
- 4: combine all $\text{split} \neq \text{split}_i$ to the training set
- 5: train S with the training set
- 6: let S predict examples in split_i
- 7: **for all** elements in prediction of S **do**
- 8: get the label associated with the highest score for the element with (3)
- 9: add the label and score to predicted points
- 10: **end for**
- 11: **end for**
- 12: sort predicted points based on the score in descending order
- 13: compute the rewards with the true label and the predicted label from the predicted points
- 14: reward points $:=$ scores from predicted points \times rewards
- 15: reduce reward points so all scores are unique
- 16: train R (optionally)
- 17: $T := \arg \max R$ or just the score which has the highest reward associated
- 18: **return** T

- Tadeusz Pietraszek. Optimizing abstaining classifiers using roc analysis. In *Proceedings of the 22Nd International Conference on Machine Learning*, ICML '05, pages 665–672, New York, NY, USA, 2005. ACM. ISBN 1-59593-180-5. doi: 10.1145/1102351.1102435. URL <http://doi.acm.org/10.1145/1102351.1102435>.
- Richard S. Sutton and Andrew G. Barto. *Reinforcement Learning: An Introduction*. MIT Press, Cambridge, MA, USA, 2nd edition, 2018.
- Francesco Tortorella. An optimal reject rule for binary classifiers. pages 611–620, 08 2000. doi: 10.1007/3-540-44522-6_63.
- Peter D. Turney. Types of cost in inductive concept learning. *CoRR*, cs.LG/0212034, 2002. URL <http://arxiv.org/abs/cs.LG/0212034>.
- Stijn Vanderlooy, Ida G. Sprinkhuizen-Kuyper, Evgueni N. Smirnov, and H. Jaap van den Herik. The roc isometrics approach to construct reliable classifiers. *Intell. Data Anal.*, 13(1):3–37, January 2009. ISSN 1088-467X. URL <http://dl.acm.org/citation.cfm?id=1551758.1551760>.

Erklärung

Ich versichere, die von mir vorgelegte Arbeit selbstständig verfasst zu haben. Alle Stellen, die wörtlich oder sinngemäß aus veröffentlichten oder nicht veröffentlichten Arbeiten anderer oder der Verfasserin/des Verfassers selbst entnommen sind, habe ich als entnommen kenntlich gemacht. Sämtliche Quellen und Hilfsmittel, die ich für die Arbeit benutzt habe, sind angegeben. Die Arbeit hat mit gleichem Inhalt bzw. in wesentlichen Teilen noch keiner anderen Prüfungsbehörde vorgelegen.

Ort, Datum

Rechtsverbindliche Unterschrift